

# CAHSI Project Report

May 23, 2024

Advisors: Professor Cihang Xie and Professor Kuan Huang

Students: Josh Angel and Oliver Chang

## 1 Motivation

Image segmentation has been in a variety of fields such as autonomous driving, medical imaging, and object-detection. Segment Anything (SA), by Meta, has further improved the vision-field by introducing the first segmentation foundation model. Segment Anything comprises of multiple promptable segmentation tasks, a transformer-based model, and data engine (providing 1+ billion masks and 11 million images). As such, Segment Anything boasts impressive zero-shot performances on novel datasets.

We seek to apply SA to medical images. In particular, we segment COVID-19 chest x-ray scans from Qatar University & Tampere University & Hamad Medical Corporation. However, computationally fine-tuning a foundation model efficiently and tailored models like ViTDet-H still outperform SA's zero-shot capabilities stand in the way from accomplishing this task. Thus we turn to Low Rank Adaptation (LoRA) to efficiently fine-tune SAM on COVID-19 lung scans.

## 2 Dataset Preparation

### 2.1 Helper functions

### 2.2 Load in subset of raw data as numpy arrays


### 2.3 Make dataset dictionary

```
[4]: {'image': <PIL.PngImagePlugin.PngImageFile image mode=RGB size=256x256>,  
      'label': <PIL.PngImagePlugin.PngImageFile image mode=L size=256x256>}
```

### 2.4 Display sample

```
[5]: (-0.5, 255.5, 255.5, -0.5)
```

Ground truth mask



## 2.5 SAM Dataset Prep

## 3 Fine-tuning

```
pixel_values torch.Size([2, 3, 1024, 1024])
original_size torch.Size([2, 2])
reshaped_input_size torch.Size([2, 2])
input_boxes torch.Size([2, 1, 4])
ground_truth_mask torch.Size([2, 256, 256])
```

SAM total params: 93729252

### 3.1 LoRA integration

#### 3.1.1 Mask Decoder

#### 3.1.2 Vision Encoder

LoRA-SAM total params: 13440996

Percentage of params reduced: 0.8565976393367569

38%| | 15/40 [00:10<00:17, 1.46it/s]

KeyboardInterrupt

Traceback (most recent call last)

Cell In[21], line 20

```
17 # Training phase
18 for batch in tqdm(train_dataloader):
19     # forward pass
--> 20     outputs = model(pixel_values=batch["pixel_values"].to(device),
21                      input_boxes=batch["input_boxes"].to(device),
22                      multimask_output=False)
23     # compute loss
24     predicted_masks = outputs.pred_masks.squeeze(1)
```

File /opt/conda/lib/python3.11/site-packages/torch/nn/modules/module.py:1501, in

```
↳ Module._call_impl(self, *args, **kwargs)
1496 # If we don't have any hooks, we want to skip the rest of the logic in
1497 # this function, and just call forward.
1498 if not (self._backward_hooks or self._backward_pre_hooks or self.
↳ _forward_hooks or self._forward_pre_hooks
1499         or _global_backward_pre_hooks or _global_backward_hooks
1500         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1501     return forward_call(*args, **kwargs)
1502 # Do not call functions when jit is used
1503 full_backward_hooks, non_full_backward_hooks = [], []
```

File /opt/conda/lib/python3.11/site-packages/transformers/models/sam/

```
↳ modeling_sam.py:1355, in SamModel.forward(self, pixel_values, input_points,
↳ input_labels, input_boxes, input_masks, image_embeddings, multimask_output,
↳ attention_similarity, target_embedding, output_attentions,
↳ output_hidden_states, return_dict, **kwargs)
1352 vision_hidden_states = None
1354 if pixel_values is not None:
-> 1355     vision_outputs = self.vision_encoder(
1356         pixel_values,
1357         output_attentions=output_attentions,
1358         output_hidden_states=output_hidden_states,
1359         return_dict=return_dict,
1360     )
1361     image_embeddings = vision_outputs[0]
1363     if output_hidden_states:
```

File /opt/conda/lib/python3.11/site-packages/torch/nn/modules/module.py:1501, in

```
↳ Module._call_impl(self, *args, **kwargs)
1496 # If we don't have any hooks, we want to skip the rest of the logic in
1497 # this function, and just call forward.
1498 if not (self._backward_hooks or self._backward_pre_hooks or self.
↳ _forward_hooks or self._forward_pre_hooks
1499         or _global_backward_pre_hooks or _global_backward_hooks
1500         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1501     return forward_call(*args, **kwargs)
1502 # Do not call functions when jit is used
```

```
1503 full_backward_hooks, non_full_backward_hooks = [], []
```

File /opt/conda/lib/python3.11/site-packages/transformers/models/sam/  
modeling\_sam.py:1043, in SamVisionEncoder.forward(self, pixel\_values,  
output\_attentions, output\_hidden\_states, return\_dict)

```
1038     layer_outputs = self._gradient_checkpointing_func(  
1039         layer_module.__call__,  
1040         hidden_states,  
1041     )  
1042 else:  
-> 1043     layer_outputs =   
-> layer_module(hidden_states, output_attentions=output_attentions)  
1045 hidden_states = layer_outputs[0]  
1047 if output_attentions:
```

File /opt/conda/lib/python3.11/site-packages/torch/nn/modules/module.py:1501, in  
Module.\_call\_impl(self, \*args, \*\*kwargs)

```
1496 # If we don't have any hooks, we want to skip the rest of the logic in  
1497 # this function, and just call forward.  
1498 if not (self._backward_hooks or self._backward_pre_hooks or self.  
-> _forward_hooks or self._forward_pre_hooks  
1499         or _global_backward_pre_hooks or _global_backward_hooks  
1500         or _global_forward_hooks or _global_forward_pre_hooks):  
-> 1501     return forward_call(*args, **kwargs)  
1502 # Do not call functions when jit is used  
1503 full_backward_hooks, non_full_backward_hooks = [], []
```

File /opt/conda/lib/python3.11/site-packages/transformers/models/sam/  
modeling\_sam.py:936, in SamVisionLayer.forward(self, hidden\_states,  
output\_attentions)

```
933     height, width = hidden_states.shape[1], hidden_states.shape[2]  
934     hidden_states, padding_shape = self.window_partition(hidden_states,  
-> self.window_size)  
--> 936 hidden_states, attn_weights = self.attn(  
937     hidden_states=hidden_states,  
938     output_attentions=output_attentions,  
939 )  
940 # Reverse window partition  
941 if self.window_size > 0:
```

File /opt/conda/lib/python3.11/site-packages/torch/nn/modules/module.py:1501, in  
Module.\_call\_impl(self, \*args, \*\*kwargs)

```
1496 # If we don't have any hooks, we want to skip the rest of the logic in  
1497 # this function, and just call forward.  
1498 if not (self._backward_hooks or self._backward_pre_hooks or self.  
-> _forward_hooks or self._forward_pre_hooks  
1499         or _global_backward_pre_hooks or _global_backward_hooks  
1500         or _global_forward_hooks or _global_forward_pre_hooks):
```

```

-> 1501     return forward_call(*args, **kwargs)
    1502 # Do not call functions when jit is used
    1503 full_backward_hooks, non_full_backward_hooks = [], []

File /opt/conda/lib/python3.11/site-packages/transformers/models/sam/
↳ modeling_sam.py:829, in SamVisionAttention.forward(self, hidden_states,
↳ output_attentions)
    826 batch_size, height, width, _ = hidden_states.shape
    827 # qkv with shape (3, batch_size, nHead, height * width, channel)
    828 qkv = (
--> 829     self.qkv(hidden_states)
    830     .reshape(batch_size, height * width, 3, self.num_attention_heads, - )
    831     .permute(2, 0, 3, 1, 4)
    832 )
    833 # q, k, v with shape (batch_size * nHead, height * width, channel)
    834 query, key, value = qkv.reshape(3, batch_size * self.
↳ num_attention_heads, height * width, -1).unbind(0)

File /opt/conda/lib/python3.11/site-packages/torch/nn/modules/module.py:1501, in
↳ Module._call_impl(self, *args, **kwargs)
    1496 # If we don't have any hooks, we want to skip the rest of the logic in
    1497 # this function, and just call forward.
    1498 if not (self._backward_hooks or self._backward_pre_hooks or self.
↳ _forward_hooks or self._forward_pre_hooks
    1499         or _global_backward_pre_hooks or _global_backward_hooks
    1500         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1501     return forward_call(*args, **kwargs)
    1502 # Do not call functions when jit is used
    1503 full_backward_hooks, non_full_backward_hooks = [], []

File /opt/conda/lib/python3.11/site-packages/loralib/layers.py:242, in
↳ MergedLinear.forward(self, x)
    240 result = F.linear(x, T(self.weight), bias=self.bias)
    241 if self.r > 0:
--> 242     result += self.lora_dropout(x) @ T(self.merge_AB().T) * self.scalin
    243 return result

File /opt/conda/lib/python3.11/site-packages/loralib/layers.py:215, in
↳ MergedLinear.merge_AB(self)
    209     return w.transpose(0, 1) if self.fan_in_fan_out else w
    210 delta_w = F.conv1d(
    211     self.lora_A.unsqueeze(0),
    212     self.lora_B.unsqueeze(-1),
    213     groups=sum(self.enable_lora)
    214 ).squeeze(0)
--> 215 return T(self.zero_pad(delta_w))

```

```

File /opt/conda/lib/python3.11/site-packages/loralib/layers.py:208, in
↳ MergedLinear.merge_AB.<locals>.T(w)
    207 def merge_AB(self):
--> 208     def T(w):
    209         return w.transpose(0, 1) if self.fan_in_fan_out else w
    210     delta_w = F.conv1d(
    211         self.lora_A.unsqueeze(0),
    212         self.lora_B.unsqueeze(-1),
    213         groups=sum(self.enable_lora)
    214     ).squeeze(0)

```

KeyboardInterrupt:

## 4 Inference

[49]:



```


pixel_values torch.Size([1, 3, 1024, 1024])
original_size torch.Size([1, 2])
reshaped_input_size torch.Size([1, 2])
input_boxes torch.Size([1, 1, 4])

(1, 1, 4)
(256, 256, 1)

```

[53]: (-0.5, 255.5, 255.5, -0.5)

Predicted mask



[54]: (-0.5, 255.5, 255.5, -0.5)

Ground truth mask



Sample IoU: 0.6381479822453883

## 5 Inference Exam

### 5.1 Load Test Set

Test data shape: (100, 256, 256, 3)

Test labels (masks) data shape: (100, 256, 256)

```
[65]: ['datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E16280_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E16628_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E16837_run-1_bp-chest_cr.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E17130_run-1_bp-chest_cr.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E17511_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E17539_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E17682_run-1_bp-chest_vp-ap_dx.png',
```



```

'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E17934_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E18100_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08047_ses-E26555_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test
Set/Images/sub-S08074_ses-E18689_run-1_bp-chest_vp-ap_cr.png']

```

```

[66]: ['datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E16280_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E16628_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E16837_run-1_bp-chest_cr.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E17130_run-1_bp-chest_cr.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E17511_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E17539_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E17682_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E17934_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E18100_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08047_ses-E26555_run-1_bp-chest_vp-ap_dx.png',
'datasets/QaTa-COV19/QaTa-COV19-v2/Test Set/Ground-
truths/mask_sub-S08074_ses-E18689_run-1_bp-chest_vp-ap_cr.png']

```

```

/opt/conda/lib/python3.11/site-packages/datasets/features/image.py:339:
UserWarning: Downcasting array dtype int64 to uint8 to be compatible with
'Pillow'

```

```

warnings.warn(f"Downcasting array dtype {dtype} to {dest_dtype} to be
compatible with 'Pillow'")

```

```

/opt/conda/lib/python3.11/site-packages/datasets/features/image.py:348:
UserWarning: Downcasting array dtype int64 to int32 to be compatible with
'Pillow'

```

```

warnings.warn(f"Downcasting array dtype {dtype} to {dest_dtype} to be
compatible with 'Pillow'")

```

```

[71]: Dataset({
    features: ['image', 'label'],
    num_rows: 100
})

```

Average IoUs over 100 test samples: 0.31679712669981375